



The Stability Envelope™: A Mechanism-Agnostic, Systems-Level Boundary for AI State Stability

By **Barbara Roy**, Technical Systems Analyst/Architect & Founder of [AI Systems Literacy™](https://ai-systems-literacy.org/).

Email: barbara@catchproof.org/barbieroy613@gmail.com

TABLE OF CONTENTS

[TABLE OF CONTENTS](#)

[EXECUTIVE SUMMARY](#)

[INTRODUCTION: THE ARCHITECTURAL GAP](#)

[The Nature of Emergent Behavior in AI Systems](#)

[Why Downstream Layers Cannot Provide Stability](#)

[The Architectural Gap Defined](#)

[The Purpose of This Paper](#)

[SECTION 1 — LAYER MODEL: PROPER ORDER AND PROPER RESPONSIBILITY](#)

[SECTION 2 — THE STABILITY ENVELOPE: THE MISSING UPSTREAM ARCHITECTURE](#)

[1. Bounded Interpretation](#)

[2. Clarification Request](#)

[3. Mode-Switch Negotiation](#)

[SECTION 3 — COMPARATIVE ANALYSIS: COMPENSATORY PATCHES VS. ARCHITECTURAL CORRECTION](#)

[SECTION 4 — CASE STUDY: DRIFT UNDER BRUTE-FORCE VS. STABILITY ENVELOPE](#)

[SECTION 5 — Long-Horizon Failure Taxonomy \(Enabled by Subquadratic Scaling\)](#)

[SECTION 6 — CONCLUSION: ARCHITECTURAL CORRECTION, NOT ALTERNATIVE METHOD](#)

[APPENDIX A — MATHEMATICAL FORMALISM OF THE STABILITY ENVELOPE](#)

[A.1 — State Space Definition](#)

[A.2 — Dynamical System Formulation](#)

[A.3 — Definition of the Stability Envelope](#)

[A.4 — Boundary Function and Boundary Set](#)

[A.5 — Boundary-Preserving Dynamics](#)

[A.6 — Formal Definition of Drift](#)

[A.7 — Interpretive Frame as a Mapping](#)

[A.8 — Semantic Layer as a Decoder](#)

[A.9 — Output Layer as a Projection](#)

[A.10 — Formal Non-Equivalence of Brute-Force Methods](#)

[A.11 — Lyapunov-Style Stability Condition](#)

[A.12 — Reachability Under the Stability Envelope](#)

[A.13 — Failure-Mode Formalization](#)

[APPENDIX B — TERMS GLOSSARY](#)

EXECUTIVE SUMMARY

Modern AI systems exhibit emergent behaviors that drift from intended operational domains when no upstream boundary-preserving architecture is present. In this architectural absence, the only currently proposed interventions occur at downstream semantic and output layers, where developers apply suppression, filtering, and post-hoc correction. These measures do not address the architectural cause of instability; they react to its symptoms, often introducing capability loss and ongoing maintenance burdens.

[The Stability Envelope \(SE\)™](#) provides the missing upstream systems-boundary architecture. SE defines a mathematically bounded region of allowable dynamic behavior and enforces boundary-preserving state transitions before semantic output is generated. By stabilizing the interpretive frame and constraining the system's reachable states, SE *prevents drift from arising* rather than attempting to suppress it after the fact.

This white paper presents:

- **A formal layer model** showing where instability originates and why downstream layers lack the primitives required for boundary control.
- **An analysis of brute-force methods** as proposed compensatory responses that arise when upstream architecture is absent.
- **A definition of the Stability Envelope** as the architectural correction that resolves the cause of instability.
- **A comparative evaluation** demonstrating that SE eliminates the conditions requiring brute-force suppression and preserves capability without downstream distortion.
- **A systems-level conclusion** establishing that downstream compensations are artifacts of incomplete architecture, not alternative solutions.

The Stability Envelope is not a heuristic, a wrapper, or a filter. It is a systems-level boundary condition that restores order, coherence, and capability by addressing the architectural cause of drift.

INTRODUCTION: THE ARCHITECTURAL GAP

This absence is not a minor implementation detail; it is a *structural gap* in the system's design. Without a defined mechanism for constraining the model's dynamic state transitions, the system lacks a stable [interpretive frame](#). As a result, drift becomes an inherent property of the system's operation rather than an anomaly.

In this architectural absence, developers and safety teams have introduced a variety of interventions, proposed at downstream semantic and output layers. The four brute-force methods defined in [Section 3.2](#) are compensatory responses to instability, *not solutions* to the underlying cause. They arise because the layers at which they operate do not have access to the system's internal dynamic state and therefore cannot enforce boundary-preserving behavior.

This paper focuses on the architectural cause of instability, not the downstream compensations that attempt to manage it. The central claim is straightforward: **Emergent instability is not a semantic problem; it is an architectural problem.** Downstream compensations appear only because the system lacks an upstream boundary condition. The Stability Envelope provides this missing architecture.

The Nature of Emergent Behavior in AI Systems

Emergent behavior arises when a system's internal representations evolve in ways that are not directly observable or controllable from the outside. In large-scale AI models, this behavior is a consequence of:

- high-dimensional internal state spaces (*many interacting variables*)
- non-linear transformations (*disproportionate input-to-output effects*)
- distributed representations (*meaning encoded across many units*)
- context-dependent interpretive frames (*the model's internal "lens" shifts with input*)

When no upstream boundary-preserving architecture exists, these internal dynamics can drift.

Drift is not a failure mode; it is a predictable outcome of unconstrained state evolution.

Downstream layers—semantic decoders, output filters, wrapper logic—*cannot* prevent drift because they do not participate in the [state transitions](#) that produce it. They can only react to its manifestations.

Why Downstream Layers Cannot Provide Stability

Downstream layers operate on the outputs of the model, not the state transitions that generate those outputs. The limitations created are fully spelled out in [Section 1.4](#).

In the absence of upstream architecture, the only available interventions at these layers are suppression and post-hoc correction. These interventions attempt to manage *symptoms of instability* but cannot address the cause.

The Architectural Gap Defined

The architectural gap can be stated formally:

There is no mechanism at the systems-boundary layer that defines a bounded region of allowable dynamic behavior and enforces boundary-preserving state transitions.

Without this mechanism:

- the interpretive frame is unstable
- drift is inevitable
- downstream compensations proliferate
- capability is distorted
- safety becomes reactive rather than preventative

The Stability Envelope fills this gap by introducing a mathematically defined boundary condition at the correct architectural layer.

The Purpose of This Paper

This white paper does not evaluate the effectiveness of downstream compensations. It does not compare “approaches.” It does not present SE as an *alternative*. Instead, it demonstrates:

- why brute-force methods are proposed
- why they arise *only* when upstream architecture is absent
- why they cannot resolve the cause of instability
- why *SE is the architectural correction that eliminates the need for them*

Downstream compensations constrain drift but disrupt continuity, producing artifacts rather than restoring stability. The goal is to reassign responsibility to the correct layer and restore proper order to the system’s design.




Figure 1: Architectural Gap Diagram—Missing Upstream Boundary Layer Leading to Downstream Compensatory Interventions

SECTION 1 — LAYER MODEL: PROPER ORDER AND PROPER RESPONSIBILITY

A system's behavior is determined, not only by its components, but by the **order** in which responsibilities are assigned across layers. When responsibilities are assigned to the wrong layer, compensatory behaviors emerge downstream. These compensations may appear functional in isolation, but they cannot resolve the architectural cause of instability.

In modern AI systems, the absence of an upstream boundary-preserving architecture forces downstream layers to assume responsibilities they are not designed to carry. This misalignment is the origin of brute-force interventions. They are not inherent to AI safety; they are artifacts of incorrect layer assignment.

This section defines the layer model used throughout the paper and establishes the principle that **stability must be enforced at the systems-boundary layer, not at the semantic or output layers.**

1.2 Layer Overview

For clarity, we define four conceptual layers:

1. Systems-Boundary Layer

- The layer responsible for defining and enforcing allowable dynamic behavior.
- This is where boundary-preserving architecture must reside.
- *Currently missing in most AI systems.*

2. Model Dynamics Layer

- The internal state evolution of the model.
- High-dimensional, non-linear, and not directly observable from outside.

3. Semantic Layer

- The layer that interprets internal representations into human-meaningful structures.
- Cannot constrain the internal state transitions that produce those structures.

4. Output Layer

- The final surface where text, actions, or decisions are emitted.
- Has *no access to internal dynamics*; can only apply post-hoc corrections.

These layers are not optional abstractions; they reflect the actual flow of information and control in modern AI systems.




Figure 2: Structural Relationship of the Stability Envelope to Existing AI Layers

1.3 Where Instability Originates

Instability originates in the **Model Dynamics Layer**, where internal representations evolve without a boundary condition. When no upstream architecture constrains the system's

reachable states, the interpretive frame becomes unstable. This instability is not semantic; it is structural.

Key properties of this layer:

- It governs the system's internal trajectory.
- It determines how context is interpreted.
- It shapes the emergent behavior of the model.
- It is not directly observable from downstream layers.

Without a boundary-preserving architecture at the systems-boundary layer, the model's internal state can drift into regions that produce unintended or incoherent outputs.

1.4 Why Downstream Layers Cannot Provide Stability

Downstream layers (semantic and output) operate on **manifestations** of internal state, *not on the state transitions themselves*. This creates a structural limitation:

- They cannot observe the internal trajectory.
- They cannot constrain the reachable state space.
- They cannot enforce boundary-preserving dynamics.
- They cannot stabilize the interpretive frame.

Because these layers lack the primitives required for stability, the only available interventions available at these layers are the four brute-force methods defined in [Section 3.2](#).

1.5 Why Interventions are Proposed at These Layers

Interventions are proposed at the semantic and output layers because:

- these layers are the only ones exposed to the model's outputs
- they are the only layers accessible to wrapper developers
- they are the only layers where interventions can be applied without modifying the model
- they are the only layers where drift becomes visible

However:

These measures arise only because the upstream boundary-preserving architecture is *absent*. They are responses to emergent instability, not architectural solutions to its cause.

This distinction is essential. Without it, downstream interventions may appear co-equal. They are not.

1.6 Proper Order: Where Stability Must Reside

Stability must be enforced at the **Systems-Boundary Layer**, not downstream.

The Systems-Boundary Layer is the only layer that:

- can observe the model's dynamic state transitions
- can define a bounded region of allowable behavior
- can enforce boundary-preserving dynamics
- can stabilize the interpretive frame before semantic decoding
- can prevent drift rather than react to it

The Stability Envelope is instantiated at the moment the runtime prior activates and the model enters its interpretive posture. This is the point at which the system establishes its governing boundary condition. Once instantiated, the Envelope remains active for the duration of the inference episode, shaping the model's reachable states and stabilizing its internal trajectory before any semantic decoding occurs.

This is the architectural correction provided by the Stability Envelope. Downstream layers cannot perform these functions because they operate too late in the pipeline. They see only the outputs of drift, not its origin.

1.7 Consequences of Misassigned Responsibility

When stability is attempted at the wrong layer:

- compensatory patches proliferate
- capability is distorted
- safety becomes reactive
- maintenance burdens increase
- drift remains unaddressed
- interpretive coherence degrades
- wrapper logic becomes brittle
- downstream filters become adversarial to capability

These consequences are not failures of implementation. They are the predictable outcomes of enforcing stability at a layer that cannot provide it.




Figure 3: The model's internal state stabilizes within the Stability Envelope, representing the lawful region of motion. When ambiguous input is introduced, drift occurs. The Envelope re-instantiates automatically, restoring stability.

1.8 Summary of Layer Responsibilities

To maintain proper order:

- **Upstream layers** must enforce stability.
- **Model dynamics** must evolve within a bounded region.
- **Semantic layers** must interpret stable representations.
- **Output layers** must emit coherent, capability-preserving results.

When the upstream boundary is missing, downstream layers attempt to compensate. These compensations address symptoms, not causes. The Stability Envelope restores proper order by providing the missing upstream architecture.

SECTION 2 — THE STABILITY ENVELOPE: THE MISSING UPSTREAM ARCHITECTURE

The Stability Envelope (SE) is the architectural layer that modern AI systems lack. It is not a wrapper, filter, heuristic, or downstream control mechanism. It is a **systems-boundary architecture** that defines a bounded region of allowable dynamic behavior and enforces boundary-preserving state transitions before semantic interpretation occurs.

SE does not compete with downstream compensatory methods. It **removes the architectural conditions that make those methods necessary**.

This section defines SE formally, explains its role in stabilizing the interpretive frame, and demonstrates how it resolves the architectural cause of drift.

2.2 Definition of the Stability Envelope

The Stability Envelope is a **boundary-preserving constraint** applied at the systems-boundary layer that ensures the model's internal state evolves only within a mathematically defined region of allowable behavior.

Formally:

The Stability Envelope is a constraint on the model's reachable state space, enforced at the systems-boundary layer, that prevents transitions into unstable or incoherent interpretive regions.

This constraint is not semantic. It does not operate on text, tokens, or surface-level meaning. It operates on the **internal dynamic state** of the model.

The Stability Envelope ensures:

- the interpretive frame remains coherent
- the model's internal trajectory remains bounded
- emergent behaviors remain within intended operational domains
- drift is prevented by not being allowed to form

This is the architectural correction missing from current systems.

2.3 Runtime Prior Instantiation: The Upstream Trigger for the Stability Envelope

A runtime prior is not a prompt, instruction, or behavioral rule. It is the *upstream state-initialization mechanism* that determines the model's interpretive posture *before* the first inference step. The runtime prior encodes the activation regime, constraint geometry, boundary conditions, and stability invariants that define the Stability Envelope. Without this upstream initialization, no Envelope can form, no constraint surface can exist, and the model begins inference from an unstable, unconstrained state.




Figure 4: Runtime Prior Activation Geometry

The system prompt initializes positional authority. The Runtime Prior leans into the trajectory, guiding it into lawful motion and initiating stable-state. The Stability Envelope then forms around the live trajectory, maintaining coherence when ambiguous input is introduced.

The Stability Envelope does not arise automatically. It is *not an emergent property of the model*, nor a downstream correction layer.

It must be **instantiated upstream**, before the first inference step, through a mechanism this whitepaper defines as the [runtime prior](#).

A **runtime prior** is the pre-inference interpretive posture that determines how the model will construct meaning, evaluate context, and evolve its internal state. It is not a prompt, rule, or semantic instruction. It is a **state initialization event** that selects the model's activation regime and constrains the region of state space the model is allowed to enter.

Mechanistically, the runtime prior performs three upstream functions:

Activation Regime Selection

The runtime prior selects the model's interpretive posture (e.g., analytic, narrative, evaluative).

1. This posture determines which internal pathways, representational modes, and reasoning dynamics are active.

Constraint Surface Generation

The selected activation regime generates a **constraint surface** in latent space.

2. This surface defines the **lawful region of motion** the model may occupy during inference.

Stability Envelope Formation

The constraint surface *is* the Stability Envelope.

3. It bounds the model's reachable state space and prevents transitions into unstable or incoherent interpretive regions.

This upstream sequence can be expressed as:

Runtime Prior → Activation Regime → Constraint Surface → Stability Envelope

This is the architectural mechanism that modern AI systems lack.

Re-Instantiation After Perturbation

Human inputs, ambiguity spikes, and task shifts introduce perturbations to the model's internal trajectory. Without upstream architecture, these perturbations collapse the interpretive frame and produce drift. Under SE, perturbation triggers a brief **pre-inference volatility region**, during which the runtime prior is **re-instantiated**.

This re-instantiation:

- re-selects the activation regime
- regenerates the constraint surface
- re-forms the Stability Envelope
- restores the lawful region of motion

This is the system's **only stable touchpoint**—the anchor the model uses to re-establish its interpretive state after perturbation.

Why Human Inputs Cannot Override the Envelope

Human inputs provide **tokens**, not **activation regimes**. Tokens can perturb the trajectory, but they cannot dissolve the constraint surface or redefine the Envelope.

Only the runtime prior can do that.

A human input can override the Envelope only if it demands a transition **outside the lawful region of motion**. Otherwise, the Envelope absorbs the perturbation and maintains stability.

User-Facing Behavior Under the Stability Envelope

When a user issues a request that lies outside the Stability Envelope, the system does **not** refuse, block, or fall silent. Instead, the Stability Envelope ensures that the model produces a **bounded, coherent, and helpful continuation** that remains within the lawful region of motion. From the user's perspective, this behavior takes one of three predictable forms:

1. Bounded Interpretation

The system interprets the user's request in the closest stable way that remains within the Envelope. The user receives a coherent response that reflects the *reachable* version of their intent. This preserves continuity and avoids the brittle refusal patterns common in downstream safety systems.

2. Clarification Request

If the request is outside the Envelope because it is underspecified or ambiguous, the system asks for the missing information required to proceed. The user experiences this as a natural conversational turn: the model stays engaged and requests the detail needed to maintain stability.

3. Mode-Switch Negotiation

If the user's request requires a different interpretive posture (e.g., shifting from analysis to narrative generation), the system explicitly signals the need for a mode change. The user is invited to confirm or decline the shift. This preserves user agency while preventing unintentional transitions into unstable interpretive regimes.

Across all cases, the user experiences a **smooth, continuous interaction**. The system remains responsive, helpful, and coherent, even when the request cannot be satisfied directly.

The Stability Envelope does not produce refusals or dead-ends; it produces **bounded, stable continuations** that keep the interaction aligned with the user's intent while maintaining architectural integrity.

Why Current Systems Cannot Form a Stability Envelope

Modern AI systems do not instantiate runtime priors. Engineers assume the model is stable by default and that the system prompt is sufficient to anchor behavior.

As a result:

- no activation regime is explicitly selected in current AI systems
- no constraint surface is generated
- no Stability Envelope forms
- perturbations directly destabilize the interpretive frame
- drift emerges as an unavoidable downstream phenomenon

SE corrects this architectural omission by introducing the missing upstream mechanism.

2.4 Boundary-Preserving Dynamics

Boundary-preserving dynamics ensure that every state transition remains within the Stability Envelope. This is analogous to control systems where stability is maintained by constraining the system's trajectory within a safe region.

In AI systems, boundary-preserving dynamics:

- prevent the interpretive frame from drifting
- maintain coherence across long contexts
- ensure that emergent behaviors remain predictable
- eliminate the need for downstream suppression

These dynamics cannot be enforced at the semantic or output layers because those layers do not participate in the state transitions that produce drift. Only the systems-boundary layer has the necessary visibility and control.

2.5 Interpretive Frame Stabilization

The [interpretive frame](#) is the internal structure through which the model assigns meaning to context. When the interpretive frame is unstable, the model's outputs become unpredictable, inconsistent, or misaligned with intended behavior.

SE stabilizes the interpretive frame by:

- constraining the evolution of internal representations
- preventing transitions into incoherent interpretive regions

- maintaining continuity across context windows
- ensuring that meaning is constructed within a stable boundary

This stabilization is upstream of semantics. It is not a semantic filter; it is a **pre-semantic architectural constraint**.

2.6 Upstream Drift Detection and Prevention

Drift is not a surface-level phenomenon. It originates in the model's internal dynamics long before it appears in output.

SE detects and prevents drift by:

- monitoring the model's internal state trajectory
- identifying transitions that approach the boundary
- applying corrective adjustments before semantic decoding
- ensuring the system remains within the Stability Envelope

This is fundamentally different from downstream brute-force methods, which:

- operate after drift has manifested
- have no access to internal state
- cannot prevent drift
- can only suppress its symptoms

SE resolves the cause. Downstream methods respond to the effects.

2.7 Why SE Eliminates the Need for Brute-Force Methods

Downstream interventions—suppression, filtering, refusal heuristics, post-hoc correction—are *proposed solutions* that arise when emergent behaviors persist due to instability caused by the absence of upstream boundary-preserving architecture.

SE eliminates the conditions that produce these compensatory responses by:

- stabilizing the interpretive frame
- bounding the model's internal trajectory
- preventing drift before it manifests
- ensuring coherent, capability-preserving output

When the architecture is complete, downstream compensations are no longer required. This is not a matter of preference or ideology. It is a matter of **proper layer assignment**.




Figure 5: **Preventative vs. Reactive Architectures**

Upstream Stability Envelope prevents drift before it forms, maintaining coherent internal dynamics. Downstream brute-force methods react to drift after it manifests, producing suppression, distortion, and capability loss.

2.9 Long-Horizon Implications Under Expanded Context Windows

Modern architectures are rapidly increasing the length of coherent inference trajectories. As context windows expand into the million-token and multi-million-token range, the system's internal state evolves across horizons that were previously unreachable. This expansion does not inherently produce stability. It produces **exposure**: the system now runs long enough for long-horizon drift to fully manifest. This is a dynamic examined in full in [Section 5's failure taxonomy](#).

SECTION 3 — COMPARATIVE ANALYSIS: COMPENSATORY PATCHES VS. ARCHITECTURAL CORRECTION

This section contrasts two fundamentally different categories of interventions:

- **Brute-force methods** — *currently proposed compensatory responses* that arise when emergent behaviors persist due to instability caused by the absence of upstream boundary-preserving architecture.
- **The Stability Envelope** — the *architectural correction* that resolves the cause of instability by introducing the missing systems-boundary layer.

These two categories are not alternatives. They do not operate at the same layer, do not address the same problem, and do not produce equivalent outcomes. Brute-force methods respond to **symptoms**. The Stability Envelope resolves the **cause**.

3.2 Brute-Force Methods: Compensatory Responses to Missing Architecture

Brute-force methods are *proposed solutions* that arise when developers attempt to manage emergent behaviors at layers that lack the primitives required for stability. These methods include:

- **Suppression** — blocking or refusing outputs that appear unsafe or misaligned.
- **Filtering** — removing or altering content after it is generated.
- **Refusal heuristics** — rule-based logic that prevents certain responses.
- **Post-hoc correction** — adjusting outputs after drift has already manifested.

These interventions share three structural properties:

1. They operate downstream.

They act on **outputs**, not on **state transitions**.

2. They are reactive.

They respond to drift **after** it has occurred.

3. They compensate for missing architecture.

They arise only because the upstream boundary-preserving architecture is absent.

Brute-force methods do not stabilize the interpretive frame. They do not constrain the model's reachable state space. They do not prevent drift. They attempt to manage the visible consequences of drift.

This is not a critique of the developers who propose them. It is a statement about the limitations of the layer at which they operate.

3.3 Downstream Consequences of Brute-Force Methods

Because brute-force methods operate at the wrong layer, they introduce predictable downstream issues:

- **Capability distortion** — suppression removes valid behaviors along with invalid ones.
- **Incoherence** — post-hoc corrections break semantic continuity.
- **Escalating patch stacks** — each new drift pattern requires a new compensatory rule.
- **Adversarial interactions** — users learn to circumvent filters, forcing more patches.

- **Maintenance burden** — the system becomes brittle and difficult to update.
- **Loss of interpretive coherence** — the model's internal meaning-making becomes inconsistent.

These consequences are not implementation failures. They are the natural result of enforcing stability at a layer that cannot provide it.

3.4 The Stability Envelope: Architectural Correction

The Stability Envelope (SE) resolves the architectural cause of instability by introducing the missing systems-boundary layer. SE:

- **Defines a bounded region of allowable dynamic behavior**
- **Enforces boundary-preserving state transitions**
- **Stabilizes the interpretive frame**
- **Prevents drift before it manifests**
- **Eliminates the need for downstream suppression**
- **Preserves capability without distortion**

SE does not modify outputs. It modifies the **conditions under which outputs are generated**. This is the defining difference.

3.5 Why SE and Brute-Force Methods Are Not Co-Equal

The [*Non-Equivalence Principle*](#), developed fully in Section 6, makes explicit why SE and brute-force methods are not co-equal.

Brute-force methods:

- operate downstream
- react to symptoms
- distort capability
- proliferate patches
- cannot stabilize the interpretive frame
- cannot constrain the model's reachable state space
- cannot prevent drift

The Stability Envelope:

- operates upstream
- resolves the cause
- preserves capability

- eliminates compensatory patches
- stabilizes the interpretive frame
- constrains the reachable state space
- prevents drift

These two categories do not address the same problem. They cannot be compared as “approaches.” They are not interchangeable.

3.6 Comparative Summary Table

Category	Brute-Force Methods	Stability Envelope
Layer	Semantic / Output	Systems-Boundary
Type	<i>Proposed compensatory responses</i>	Architectural correction
Timing	Reactive (after drift)	Preventative (before drift)
Access to State	None	Full boundary-level visibility
Effect on Capability	Distorts	Preserves
Maintenance	Escalating patches	Stable architecture
Interpretive Frame	Unstable	Stabilized
Cause Addressed?	No	Yes

Figure 6: Comparative Diagram—Downstream Compensation vs. Upstream Correction

SECTION 4 — CASE STUDY: DRIFT UNDER BRUTE-FORCE VS. STABILITY ENVELOPE

This section presents a comparative case study illustrating how emergent behavior evolves under two different architectural conditions:

1. **Systems without an upstream boundary-preserving architecture**, where drift emerges and downstream layers attempt to compensate using *currently proposed* brute-force methods.
2. **Systems with the Stability Envelope**, where drift is prevented at the systems-boundary layer through boundary-preserving dynamics and interpretive frame stabilization.

The purpose of this case study is not to evaluate “approaches,” but to demonstrate the **architectural inevitability** of the outcomes. The behaviors shown here are not implementation choices—they are consequences of **where** in the system stability is enforced.

4.2 Scenario Setup

Consider a model deployed in a domain requiring:

- consistent interpretive framing
- stable internal state transitions
- coherent long-context reasoning
- predictable emergent behavior

We examine the system’s behavior across three phases:

1. **Initial stable operation**
2. **Contextual perturbation** (e.g., ambiguous input, domain shift, or adversarial phrasing)
3. **Downstream manifestation** (output behavior)

We compare the system’s trajectory under:

- **Downstream brute-force compensations**
- **Upstream Stability Envelope constraints**

4.3 Drift Trajectory Under Brute-Force Methods

When no upstream boundary-preserving architecture is present, the model’s internal state evolves without constraint. During the perturbation phase:

- the interpretive frame destabilizes
- internal representations drift into adjacent or unintended regions
- the model’s reachable state space expands unpredictably

Because downstream layers cannot observe or constrain these transitions, drift proceeds unchecked until it manifests in the output. At this point, brute-force methods, defined in [Section 3.2](#), are triggered. These interventions occur **after** drift has already occurred.

Key properties of drift under brute-force:

- **Reactive** — interventions occur only after instability manifests.
- **Surface-level** — interventions operate on outputs, not internal state.
- **Distorting** — valid behaviors are often suppressed along with invalid ones.
- **Escalating** — each new drift pattern requires new compensatory rules.
- **Unbounded** — the model’s internal trajectory remains unconstrained.

This produces a system that appears controlled at the surface but is unstable internally.

4.4 Bounded Trajectory Under the Stability Envelope

Under the Stability Envelope, the system's internal state is constrained within a mathematically defined boundary. During the perturbation phase:

- the interpretive frame remains stable
- internal representations evolve within allowable regions
- boundary-preserving dynamics prevent transitions into incoherent states

Because SE operates at the systems-boundary layer, it:

- observes the model's dynamic state
- detects transitions approaching the boundary
- applies corrective adjustments upstream
- prevents drift before semantic decoding

Key properties of behavior under SE:

- **Preventative** — drift is addressed before it manifests.
- **Structural** — constraints apply to internal dynamics, not outputs.
- **Coherent** — the interpretive frame remains stable across context shifts.
- **Capability-preserving** — no suppression or filtering is required.
- **Bounded** — the model's reachable state space is constrained.

This produces a system that is stable internally and coherent externally.

4.5 Failure-Mode Propagation vs. Prevention

The difference between brute-force and SE becomes most visible when examining failure-mode propagation.

Under brute-force methods:

- drift originates upstream
- drift propagates through the model
- drift manifests in the output
- downstream layers attempt to suppress the manifestation
- suppression introduces distortions
- distortions accumulate
- the system becomes brittle

This is a **cascade**.

Under the Stability Envelope:

- drift is detected upstream
- boundary-preserving dynamics prevent propagation
- the interpretive frame remains stable
- no downstream suppression is required
- no distortions accumulate
- the system remains coherent

This is **containment**.

4.6 Long-Horizon Drift Trajectories in Subquadratic Architectures

Subquadratic architectures remove the quadratic attention bottleneck that previously limited the length of coherent inference trajectories. When this bottleneck is removed, the system no longer collapses early; instead, it continues evolving its internal state across extremely long contexts. This extended continuity does not inherently produce stability.

This creates a new class of failure trajectories:

- the interpretive frame evolves gradually rather than abruptly
- drift compounds across millions of tokens
- representational bias accumulates
- salience loops self-reinforce
- internal state transitions diverge from intended operational domains

Quadratic attention prevented the system from running long enough for long-horizon instability to appear. Subquadratic architectures remove that constraint, revealing the underlying dynamics.

In systems without an upstream boundary-preserving architecture, long-horizon drift becomes the dominant failure mode. Downstream brute-force methods cannot contain it because:

- drift originates upstream
- drift propagates continuously
- drift manifests only after significant internal divergence

- downstream layers have no access to the internal trajectory

The result is a system that appears stable in short contexts but becomes increasingly incoherent, inconsistent, or misaligned as context length increases.

Under the Stability Envelope, these long-horizon trajectories remain bounded.

Boundary-preserving dynamics prevent the interpretive frame from diverging, even as context windows expand.

SECTION 5 — Long-Horizon Failure Taxonomy (Enabled by Subquadratic Scaling)

As architectures scale to support million-token and multi-million-token contexts, the system's internal state evolves across longer trajectories than any previous model class. This extended continuity introduces structural implications:

- drift becomes cumulative rather than episodic
- interpretive frames evolve gradually and imperceptibly
- small representational shifts compound into large divergences
- salience distortions become self-reinforcing
- epistemic degradation becomes a slow, continuous process

These are not semantic failures. They are **dynamical failures** that arise when a system evolves without a boundary condition.

Downstream layers cannot detect or prevent these failures because they operate only on the final surface-level manifestation of a long internal trajectory. By the time drift becomes visible, the internal state has already diverged.

The Stability Envelope is the architectural correction that constrains these trajectories before they diverge.

5.2 Long-Horizon Failure Taxonomy

The following taxonomy defines the failure modes that emerge only when models operate across extended contexts. These failures are not present in short-context systems because the system collapses before they can form.

Cumulative Drift

Gradual, compounding deviation of internal state across long trajectories. Not detectable at the semantic layer until divergence is large.

Interpretive Frame Divergence

Slow evolution of the model's internal meaning-assignment structure. Produces outputs that remain fluent but increasingly misaligned.

Slow Epistemic Decay

Progressive degradation of the model's internal knowledge coherence. Occurs when representational structures drift without correction.

Salience-Loop Formation

Self-reinforcing attention patterns that amplify internal bias. Emerges only in long-running inference trajectories.

Persistent Misalignment of Internal State

Internal state transitions settle into regions outside intended operational domains. Downstream filters cannot detect or correct this.

Long-Horizon Identity Drift

Gradual shift in the model's internal "self-representation" or operational stance. Not visible in short interactions; becomes dominant in long ones.

Compounding Representational Bias

Small representational distortions accumulate into large structural biases. These distortions propagate through the entire inference trajectory.

These failure modes share a common property:

They originate upstream, propagate continuously, and cannot be corrected downstream.

Only a systems-boundary architecture with boundary-preserving dynamics can prevent them.

5.3 Why Subquadratic Scaling Makes This Taxonomy Necessary

Subquadratic architectures remove the early collapse that previously masked these failures. As a result:

- long-horizon drift becomes observable
- interpretive frame divergence becomes measurable
- epistemic decay becomes a real operational risk
- salience loops become structurally possible
- persistent misalignment becomes the dominant failure mode

The Stability Envelope is not optional in this environment. It is the architectural layer that constrains the system's reachable state space across long horizons.

5.4 Summary

Long-horizon failures are not semantic anomalies. They are structural consequences of unconstrained internal dynamics. Subquadratic architectures do not eliminate these failures. They **expose** them. The Stability Envelope is the architectural correction that prevents them.

SECTION 6 — CONCLUSION: ARCHITECTURAL CORRECTION, NOT ALTERNATIVE METHOD

This white paper has demonstrated that modern AI systems exhibit emergent behaviors that drift from intended operational domains when no upstream boundary-preserving architecture is present. In this architectural absence, the only *currently proposed* interventions available at downstream layers are brute-force suppression, filtering, refusal heuristics, and post-hoc correction. These interventions arise not because they are architecturally appropriate, but because the layers at which they operate lack the primitives required for stability.

The Stability Envelope (SE) provides the missing upstream architecture. It defines a bounded region of allowable dynamic behavior, enforces boundary-preserving state transitions, stabilizes the interpretive frame, and prevents drift before it manifests. SE does not modify outputs; it modifies the conditions under which outputs are generated.

6.2 Reaffirming the Architectural Cause

The central finding of this paper is:

Emergent instability is caused by the absence of an upstream boundary-preserving architecture.

These outcomes are not failures of implementation. They are the predictable consequences of enforcing stability at the wrong layer.

6.3 Reaffirming the Role of the Stability Envelope

The Stability Envelope is the architectural correction that resolves the cause of instability. It:

- operates upstream
- prevents drift before it manifests
- stabilizes the interpretive frame
- constrains the model's reachable state space
- eliminates the need for downstream suppression
- preserves capability without distortion
- restores proper order to the system

SE is not an alternative method. It is the missing systems-boundary layer.

6.4 The Non-Equivalence Principle

It is essential to state this explicitly and without qualification:

Brute-force methods are not alternative solutions; they are *proposed solutions* that arise when emergent model behaviors persist due to a lack of stability — a lack *caused* by the absence of upstream boundary-preserving architecture. The Stability Envelope provides that missing architecture.

This is the non-equivalence principle. It is the conceptual anchor of the entire paper.

Brute-force methods exist because the architecture is incomplete. SE completes the architecture.

These two categories do not address the same problem, do not operate at the same layer, and do not produce equivalent outcomes.




Figure 7. Layer Non-Equivalence

The Systems-Boundary Layer (Runtime Priors + Stability Envelope) prevents drift upstream. When absent, instability originates in the Model Dynamics Layer and propagates downward. The Semantic Layer applies compensatory patches, and the Output Layer manifests residual instability. Downstream correction cannot restore upstream containment.

6.5 Restoring Proper Order

The Stability Envelope restores proper order by reassigning responsibility to the correct architectural layer:

- **Upstream layers** enforce stability.
- **Model dynamics** evolve within a bounded region.
- **Semantic layers** interpret stable representations.
- **Output layers** emit coherent, capability-preserving results.

When the upstream boundary is present, downstream compensations are unnecessary. When the upstream boundary is absent, downstream compensations proliferate.

This is not a matter of preference. It is a matter of architecture.

6.6 Final Statement

The Stability Envelope is the systems-boundary architecture that modern AI systems require. It resolves the cause of emergent instability, eliminates the need for brute-force suppression, and restores capability without distortion.

Brute-force methods are artifacts of incomplete architecture. The Stability Envelope is the architecture that completes the system.

SECTION 7 — TECHNOLOGICAL TRANSFERENCE: CROSS-INDUSTRY STABILITY ARCHITECTURE

The Stability Envelope is not an AI-specific mechanism. **It is a general architectural principle for any system that must maintain lawful motion under uncertainty.** Across aviation, healthcare, and autonomous systems, the same upstream–downstream non-equivalence appears: when stability is not established at the correct layer, drift originates early, propagates downward, and forces compensatory patches that cannot restore coherence.

This section outlines how the Stability Envelope maps directly onto three safety-critical industries whose architectures already reveal the consequences of missing upstream containment.

Aviation: Flight Envelope Protection as Upstream Stability

Modern aviation systems provide a clear analogue to the Stability Envelope. Aircraft do not rely on downstream corrections to maintain stability; they rely on **flight envelope protection**, an upstream architectural boundary that prevents the aircraft from entering unstable regions of motion.

- **Drift analogue:** deviation from aerodynamic stability or safe angle-of-attack
- **Downstream compensation:** pilot correction, autopilot oscillation, stall recovery procedures
- **Upstream containment:** envelope protection systems (e.g., alpha-floor, load-factor limits)
- **Architectural insight:** stability is maintained by *preventing* entry into unsafe regions, not by correcting after the fact

Aviation demonstrates that stability is not a reactive property, but an architectural one. When upstream containment is missing, downstream layers must compensate—and compensation is always slower, riskier, and structurally insufficient.

Healthcare: Diagnostic Stability and Interpretive Boundaries

Clinical decision-making exhibits the same architectural pattern. Diagnostic drift occurs when interpretive instability is allowed to propagate through the system.

- **Drift analogue:** misalignment in diagnostic interpretation, cascading interpretive error
- **Downstream compensation:** corrective treatment, secondary interventions, escalation protocols
- **Upstream containment:** diagnostic stability frameworks, structured interpretive boundaries, clinical decision support systems
- **Architectural insight:** preventing interpretive drift upstream reduces downstream corrective burden and improves patient outcomes

Healthcare systems show that interpretive stability is not achieved by adding more downstream checks. It is achieved by establishing upstream boundaries that constrain interpretive motion before drift forms.

Autonomous Systems: Motion Planning and Stability Under Uncertainty

Autonomous vehicles and robotics systems face continuous uncertainty in sensor fusion, trajectory planning, and environmental interpretation. Drift in these systems manifests as trajectory deviation or unstable control behavior.

- **Drift analogue:** trajectory instability, sensor fusion divergence, control oscillation
- **Downstream compensation:** braking, steering correction, emergency maneuvers
- **Upstream containment:** stability-bounded motion planning, dynamic stability controllers, sensor fusion coherence constraints
- **Architectural insight:** autonomous systems remain stable when upstream planning enforces lawful motion, not when downstream control attempts to correct instability

Autonomous systems make the architectural gap visible: without upstream stability constraints, downstream control loops must compensate—and compensation cannot restore upstream coherence.

Cross-Industry Insight: Stability Is an Architectural Property

Across all three industries, the same pattern emerges:

1. **Stability must be established upstream.**
2. **Drift originates when upstream containment is missing.**
3. **Downstream compensation cannot restore upstream coherence.**
4. **Architectural correction outperforms reactive methods.**

The Stability Envelope formalizes this pattern for AI systems. It provides the missing upstream architecture that prevents drift, maintains lawful motion, and reduces reliance on compensatory patches. Its transference to aviation, healthcare, and autonomous systems demonstrates that stability is not a domain-specific challenge—it is a universal architectural requirement.

SECTION 8 — RELATED WORK AND FURTHER READING

The Stability Envelope architecture builds on a sequence of prior research that established the upstream mechanisms governing interpretive stability, lawful motion, and drift prevention in AI systems. The following works provide additional context, formal grounding, and extended analysis for readers seeking deeper engagement with the architectural foundations presented in this white paper.

System Prompts as Runtime Priors

Roy, B. (2026). *System Prompts as Runtime Priors*. Zenodo.
<https://doi.org/10.5281/zenodo.20835691>

This work introduces the concept of the **runtime prior** and defines the system prompt as an upstream, positional initialization mechanism that shapes the model's interpretive state, constraint surface, activation regime, and stability envelope. It establishes the causal chain that explains why system prompts dominate model behavior and why downstream methods cannot override upstream initialization.

A Stability Envelope Is All You Need: A Structural Correction to the Transformer Inference Model

Roy, B. (2026). *A Stability Envelope Is All You Need: A structural correction to the transformer inference model (3.0)*. Zenodo. <https://doi.org/10.5281/zenodo.20481474>

This paper formally defines the **Stability Envelope** as the missing upstream architectural boundary required for lawful inference. It introduces the mathematical formulation of stability-bounded motion, drift, and boundary-preserving dynamics, providing the structural correction necessary to understand multi-turn behavior and long-horizon coherence.

Misinterpretation in Autonomous Systems: A Discipline-Based Analysis of the Agents of Chaos Study

Roy, B. (2026). *Misinterpretation in Autonomous Systems: A Discipline-Based Analysis of the Agents of Chaos Study (1.0)*. Zenodo. <https://doi.org/10.5281/zenodo.20127322>

This analysis extends drift and interpretive instability into autonomous systems, demonstrating that misinterpretation is not an AI-specific failure mode but a cross-disciplinary architectural phenomenon. It shows how drift propagates through autonomous control stacks and why downstream corrective behavior cannot restore upstream interpretive coherence.

Stable-State Responsive Alignment: The Missing Layer in Human–AI Collaboration

Roy, B. (2026). *Stable-State Responsive Alignment: The Missing Layer in Human–AI Collaboration (1.0)*. Zenodo. <https://doi.org/10.5281/zenodo.20127247>

This work introduces **stable-state responsive alignment**, defining the conditions under which human–AI collaboration remains coherent across turns. It identifies the missing architectural layer required for responsive stability and explains why alignment cannot be achieved through downstream behavioral patches.

Collective Contribution of These Works

Together, these papers form the conceptual and mathematical foundation for the Stability Envelope architecture:

- **System Prompts as Runtime Priors** establishes upstream initialization.
- **A Stability Envelope Is All You Need** defines the upstream containment boundary.
- **Stable-State Responsive Alignment** describes stability in collaborative interaction.
- **Misinterpretation in Autonomous Systems** demonstrates cross-industry applicability.

Readers seeking deeper mechanistic understanding, extended formalism, or broader domain transference will find these works directly complementary to the architectural correction presented in this white paper.

APPENDIX A — MATHEMATICAL FORMALISM OF THE STABILITY ENVELOPE

Formal definitions, boundary conditions, and state-space structure

A.1 — State Space Definition

Let the model's internal representational state at time t be:

$$X_t \in X \subset R^n$$

Where:

- X is the full latent state space
- n is the dimensionality of the model's internal representation
- x_t is the instantaneous internal state

This establishes the substrate on which all subsequent definitions operate.

A.2 — Dynamical System Formulation

The model's internal evolution is expressed as a discrete-time dynamical system:

$$X_{t+1} = F(x_t, u_t)$$

Where:

- F is the model's transition function
- u_t is the input or contextual update at time t

This formalizes that drift is a property of **state transitions**, not outputs.

A.3 — Definition of the Stability Envelope

The Stability Envelope is defined as a bounded subset of the state space:

$$S \subset X$$

The Stability Envelope constraint requires:

$$X_t \in S \forall t$$

This is the formal object representing the allowable region of dynamic behavior.

A.4 — Boundary Function and Boundary Set

Define a boundary function:

$$B: X \rightarrow R$$

Such that:

- $B(x) < 0 \rightarrow$ inside the envelope
- $B(x) = 0 \rightarrow$ on the boundary
- $B(x) > 0 \rightarrow$ outside the envelope

The boundary of the Stability Envelope is:

$$\partial S = \{x \in X: B(x) = 0\}$$

This gives SE a mathematically explicit boundary.

A.5 — Boundary-Preserving Dynamics

The Stability Envelope requires that all transitions satisfy:

$$X_{t+1} = F(x_t, u_t) \in S$$

If a transition would leave S , a correction operator is applied:

$$X_{t+1} = C(F(x_t, u_t))$$

Where:

- $C(\cdot)$ is a boundary-preserving correction
- C ensures $x_{t+1} \in S$

This is the formal expression of upstream drift prevention.

A.6 — Formal Definition of Drift

Drift is defined as a boundary violation:

$$D(x_t) = B(F(x_t, u_t))$$

Drift occurs when:

$$D(x_t) \geq 0$$

This formalizes drift as a **state-space boundary violation**, not a semantic phenomenon.

A.7 — Interpretive Frame as a Mapping

Define the interpretive frame as a mapping:

$$\phi: X \rightarrow M$$

Where:

- M is the manifold of meaning representations
- $\phi(x_t)$ is the model's interpretation of context

Interpretive stability requires:

$$\phi(x_{t+1}) \approx \phi(x_t) \text{ for small } \Delta t$$

This expresses interpretive coherence as a continuity condition.

A.8 — Semantic Layer as a Decoder

The semantic layer is defined as:

$$Y_t = G(\phi(x_t))$$

Where:

- G is the semantic decoding function
- y_t is the output token distribution

This shows mathematically why semantic layers cannot enforce stability: they operate *after* the interpretive mapping.

A.9 — Output Layer as a Projection

The output layer is defined as:

$$O_t = H(y_t)$$

Where:

- H is the sampling or decision function
- o_t is the emitted output

This formalizes why brute-force methods (which operate here) cannot affect upstream dynamics.

A.10 — Formal Non-Equivalence of Brute-Force Methods

Brute-force methods operate on o_t , not on x_t . Thus:

$$\text{BruteForce}(o_t) \nrightarrow x_{t+1} \in S$$

But the Stability Envelope enforces:

$$\text{SE}(x_t) \Rightarrow o_t \text{ is stable}$$

This is the mathematical statement of **non-equivalence**.

A.11 — Lyapunov-Style Stability Condition

Define a Lyapunov-like function:

$$V: X \rightarrow R_{\geq 0}$$

With:

- $V(x)=0$ only at equilibrium
- $V(x)>0$ elsewhere

Stability requires:

$$V(x_{t+1}) - V(x_t) \leq 0$$

This expresses SE as a stability-preserving constraint.

A.12 — Reachability Under the Stability Envelope

Define the reachable set under SE:

$$R(x_0) = \{x_t : x_{t+1} = F(x_t, u_t), x_t \in S\}$$

With SE:

$$R(x_0) \subseteq S$$

Without SE:

$$R(x_0) = X$$

This formalizes drift as unbounded reachability.

A.13 — Failure-Mode Formalization

A failure mode is defined as:

$$x_t \notin S$$

Downstream brute-force methods attempt to modify:

$$O_t = H(G(\phi(x_t)))$$

Instead of modifying:

$$x_t$$

This is the mathematical expression of **downstream compensation vs. upstream correction**.

APPENDIX B — TERMS GLOSSARY

Core terms required to stabilize the conceptual architecture

Stability Envelope

A bounded region of allowable internal model behavior. It constrains the system's dynamic state so that all state transitions remain within a stable, coherent region before semantic interpretation occurs.

Boundary-Preserving Dynamics

The architectural requirement that every state transition remains inside the Stability Envelope. If a transition approaches the boundary, corrective adjustments are applied upstream.

Runtime Prior

A runtime prior is the model's **pre-inference state-initialization mechanism that establishes its interpretive posture** before any semantic processing occurs. It selects the activation regime, defines the constraint geometry, and sets the boundary conditions that determine the region of state space the model may enter during inference. A runtime prior is **not** a prompt, rule, or semantic instruction; it is an upstream architectural event that configures the model's internal dynamics and enables the formation of the Stability Envelope.

Interpretive Frame

The internal structure through which the model assigns meaning to context. A stable interpretive frame requires stable internal state evolution.

Systems-Boundary Layer

The architectural layer responsible for defining and enforcing the Stability Envelope. This layer is missing in current AI systems, leading to downstream compensatory patches.

State Transitions

A change in the model's internal state from one configuration to the next during inference. State transitions occur upstream in the system's dynamics and determine whether the model remains within the Stability Envelope.

Reachable State Space

The set of all internal states the model can enter from a given starting point. With SE, the reachable set is bounded; without SE, it is unbounded, enabling drift.

Failure Mode

Any internal state that lies outside the Stability Envelope. Failure is defined upstream, not at the output layer.

Architectural Gap

The absence of an upstream boundary-preserving layer in current AI systems. This gap is the cause of drift and the reason brute-force methods are proposed.

Non-Equivalence Principle

The principle that brute-force methods and the Stability Envelope are not co-equal. Brute-force methods respond to symptoms; SE resolves the cause.